



Reactive Dashboards Using Apache Spark

Rahul Kumar

Software Developer


@rahul_kumar_aws

LinuxCon, CloudOpen, ContainerCon North America 2015

Agenda



- Big Data Introduction
- Apache Spark
- Introduction to Reactive Applications
- Reactive Platform
- Live Demo



A typical database application

Impression

49,576,760

Multi
Source
Data
Ingestion

Gb's to
Petabyte
Data

Realtime
update

Conversion

670,304

Metrics

Download

Impression



Sub
second
response

Name

Impression


CH	5003590
CN	4945360
FR	4946540
HK	4942300
IN	4965480
JP	4898860
PK	4984810
SP	4961700
US	4961700

Impression

chrome	9933570
firefox	9933570
ie	9933570
opera	9933570
safari	9933570

Scalable

Click



Device

Name

Impression


Connected Device	8330540
Connected TV	8317610
Mobile	8224120
Personal Computer	8173370
Set Top Box	8296440
Tablet	8234680

Site


Name

Impression


amazon.in	4987980
fb.com	4944110
flipkart.com	4939570
google.com	4977100
jabong.com	4956230
myntra.com	4938360
quora.com	4939470
snapdeal.com	4957240
timeofindia.com	4957240



Three V's of Big Data



Scale vertically (scale up)



Scale horizontally (scale out)

Apache Spark



Apache Spark is a fast and general engine for large-scale data processing.

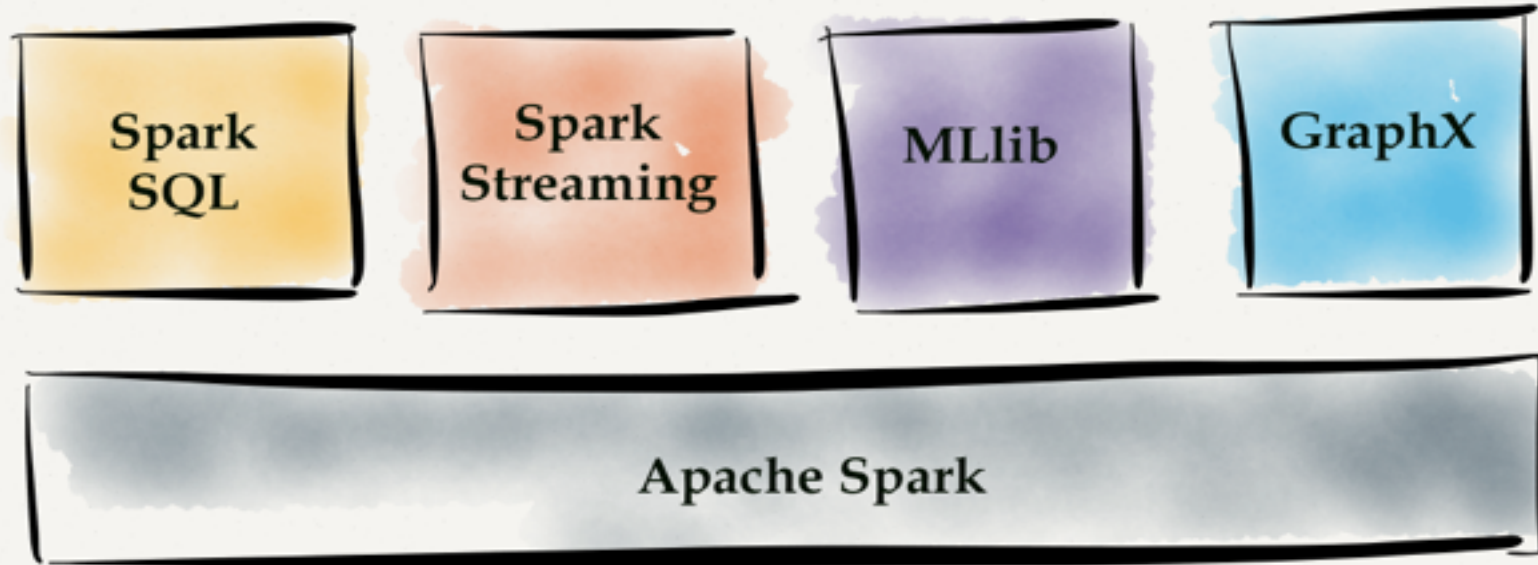
Speed

**Easy to
Use**

Generality

**Runs
Everywhere**

Apache *Spark* Stack



- Apache Spark Setup
- Interaction with Spark Shell
- Setup a Spark App
- RDD Introduction
- Deploy Spark app on Cluster

Prerequisite for cluster setup



Spark Cluster

Spark runs on Java 6+, Python 2.6+ and R 3.1+.

For the Scala API, Spark 1.4.1 uses Scala 2.10.

Java 8

```
sudo add-apt-repository ppa:webupd8team/java
```

```
$ sudo apt-get update
```

```
$ sudo apt-get install oracle-java8-installer
```

Scala 1.10.4

```
http://www.scala-lang.org/files/archive/scala-2.10.4.tgz
```

```
$tar -xvzf scala-2.10.4.tgz
```

```
vim ~/.bashrc
```

```
export SCALA_HOME=/home/ubuntu/scala-2.10.4
```

```
export PATH=$PATH:$SCALA_HOME/bin
```

Spark Setup



<http://spark.apache.org/downloads.html>

Download Spark

The latest release of Spark is Spark 1.4.1, released on July 15, 2015 ([release notes](#)) ([git tag](#))

1. Choose a Spark release ✓ 1.4.1 (Jul 15 2015)
1.4.0 (Jun 11 2015)
1.3.1 (Apr 17 2015)
1.3.0 (Mar 13 2015)
1.2.2 (Apr 17 2015)
1.2.1 (Feb 09 2015)
1.2.0 (Dec 18 2014)
1.1.1 (Nov 26 2014)
1.1.0 (Sep 11 2014)
1.0.2 (Aug 05 2014)
0.9.2 (Jul 23 2014)
0.8.1 (Dec 19 2013)
0.8.0 (Sep 25 2013)
0.7.3 (Jul 16 2013)
2. Choose a package type
3. Choose a download type
4. Download Spark: [spark](#)
5. Verify this release using

Note: Scala 2.11 users should

Link with Spark


several Hadoop versions]

id checksums.

source package and build *with Scala 2.11 support*.

Download Spark

The latest release of Spark is Spark 1.4.1, released on July 15, 2015 ([release notes](#)) ([git tag](#))

1. Choose a Spark release: 1.4.1 (Jul 15 2015) 
2. Choose a package type: ☒ Source Code [can build several Hadoop versions]
Pre-built with user-provided Hadoop [can use with most Hadoop distributions]
3. Choose a download type: ☐ Pre-built for Hadoop 2.6 and later
☐ Pre-built for Hadoop 2.4 and later
☐ Pre-built for Hadoop 2.3
☐ Pre-built for Hadoop 1.X
☐ Pre-built for CDH 4
4. Download Spark: [spark](#)
5. Verify this release using

Note: Scala 2.11 users should download the Spark source package and build [with Scala 2.11 support](#).


Running Spark Example & Shell



```
$ cd spark-1.4.1-bin-hadoop2.6
```

```
$/bin/run-example SparkPi 10
```

```
15/08/08 21:26:16 INFO TaskSetManager: Finished task 6.0 in stage 0.0 (TID 6) in 55 ms on localhost (9/10)
15/08/08 21:26:16 INFO TaskSetManager: Finished task 7.0 in stage 0.0 (TID 7) in 92 ms on localhost (10/10)
15/08/08 21:26:16 INFO DAGScheduler: ResultStage 0 (reduce at SparkPi.scala:35) finished in 0.829 s
15/08/08 21:26:16 INFO TaskSchedulerImpl: Removed TaskSet 0.0, whose tasks have all completed, from pool
15/08/08 21:26:16 INFO DAGScheduler: Job 0 finished: reduce at SparkPi.scala:35, took 1.110582 s
Pi is roughly 3.144248
15/08/08 21:26:16 INFO SparkUI: Stopped Spark web UI at http://192.168.1.117:4040
15/08/08 21:26:16 INFO DAGScheduler: Stopping DAGScheduler
15/08/08 21:26:16 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
15/08/08 21:26:16 INFO Utils: path = /private/var/folders/c5/d743br192cy_kh_26ns38k7m0000ap/T/spark-ad5e0ddf-ad19-d
```



```

Welcome to
general spark-submit script for launching applications). For example,

  /run-example SparkPi 10

version 1.4.1

Using Scala version 2.10.4 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_45)
Type in expressions to have them evaluated.
Type :help for more information.
15/08/08 21:42:09 INFO SparkContext: Running Spark version 1.4.1
15/08/08 21:42:09 INFO SecurityManager: Changing view acls to: rahul
15/08/08 21:42:09 INFO SecurityManager: Changing modify acls to: rahul
15/08/08 21:42:17 INFO HiveMetaStore: Added admin role in metastore
15/08/08 21:42:17 INFO HiveMetaStore: Added public role in metastore
15/08/08 21:42:18 INFO HiveMetaStore: No user is added in admin role, since config is empty
15/08/08 21:42:18 INFO SessionState: No Tez session required at this point. hive.execution.engine=mr.
15/08/08 21:42:18 INFO SparkILoop: Created sql context (with Hive support)...
SQL context available as sqlContext.

scala>

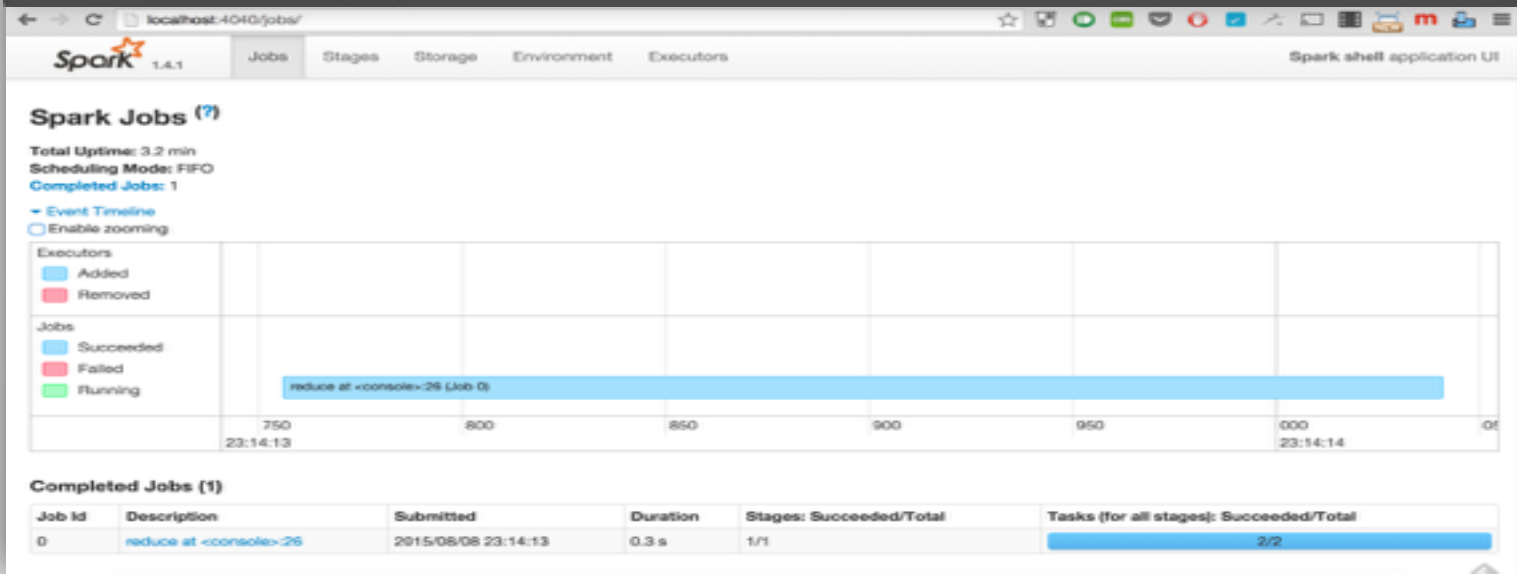
```

The `--master` option specifies the master URL for a distributed cluster, or `local` to run locally with one thread, or `local[N]` to run locally with N threads.


```
scala> val data = Array(1, 2, 3, 4, 5)
data: Array[Int] = Array(1, 2, 3, 4, 5)
```

```
scala> val distData = sc.parallelize(data)
distData: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[0] at parallelize at <console>:23
```

```
scala> distData.reduce(_+_)
```



RDD Introduction

Resilient Distributed Data Set

Resilient Distributed Datasets (RDDs), a ***distributed memory abstraction*** that lets programmers perform ***in-memory computations*** on large clusters in a ***fault-tolerant*** manner.

RDD shard the data over a cluster, like a virtualized, distributed collection.


Users create **RDDs** in two ways: by **loading an external dataset**, or by **distributing a collection of objects** such as List, Map etc.

RDD Operations

RDDs support two types of operations: **transformations** and **actions**.

Spark computes RDD only in a **lazy fashion**.

Only computation start when an **Action** call on RDD.



- Simple SBT project setup <https://github.com/rahulkumar-aws/HelloWorld>



```
$ mkdir HelloWorld
$ cd HelloWorld
$ mkdir -p src/main/scala
$ mkdir -p src/main/resources
$ mkdir -p src/test/scala
$ vim build.sbt
    name := "HelloWorld"

    version := "1.0"

    scalaVersion := "2.10.4"
$ mkdir project
$ cd project
$ vim build.properties
    sbt.version=0.13.8

$ vim src/main/scala/HelloWorld.scala

object HelloWorld { def main(args: Array[String]) = println("HelloWorld!") }
$ sbt run
```

First Spark Application



\$git clone <https://github.com/rahulkumar-aws/WordCount.git>

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._
object SparkWordCount {

    def main(args: Array[String]): Unit = {

        val sc = new SparkContext("local", "SparkWordCount")

        val wordsCounted = sc.textFile(args(0)).map(line=> line.toLowerCase)
                                .flatMap(line => line.split("\\W+"))
                                .groupBy(word => word)
                                .map{ case(word, group) => (word, group.size)}


        wordsCounted.saveAsTextFile(args(1))
        sc.stop()
    }
}
```

\$sbt "run-main ScalaWordCount src/main/resources/sherlockholmes.txt out"

Launching Spark on Cluster



Spark cluster components



Cluster Manager Can be Spark's own Standalone Cluster Manager or Mesos or YARN

Spark Cache Introduction



Spark supports pulling data sets into a cluster-wide in-memory cache.

```
scala> val textFile = sc.textFile("README.md")

textFile: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[12] at textFile at <console>:21

scala> val linesWithSpark = textFile.filter(line => line.contains("Spark"))

linesWithSpark: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[13] at filter at
<console>:23

scala> linesWithSpark.cache()

res11: linesWithSpark.type = MapPartitionsRDD[13] at filter at <console>:23

scala> linesWithSpark.count()

res12: Long = 19
```

← → C localhost:4040/storage/ ☆ [Icons]

Spark 1.4.1 Jobs Stages **Storage** Environment Executors Spark shell application UI

Storage

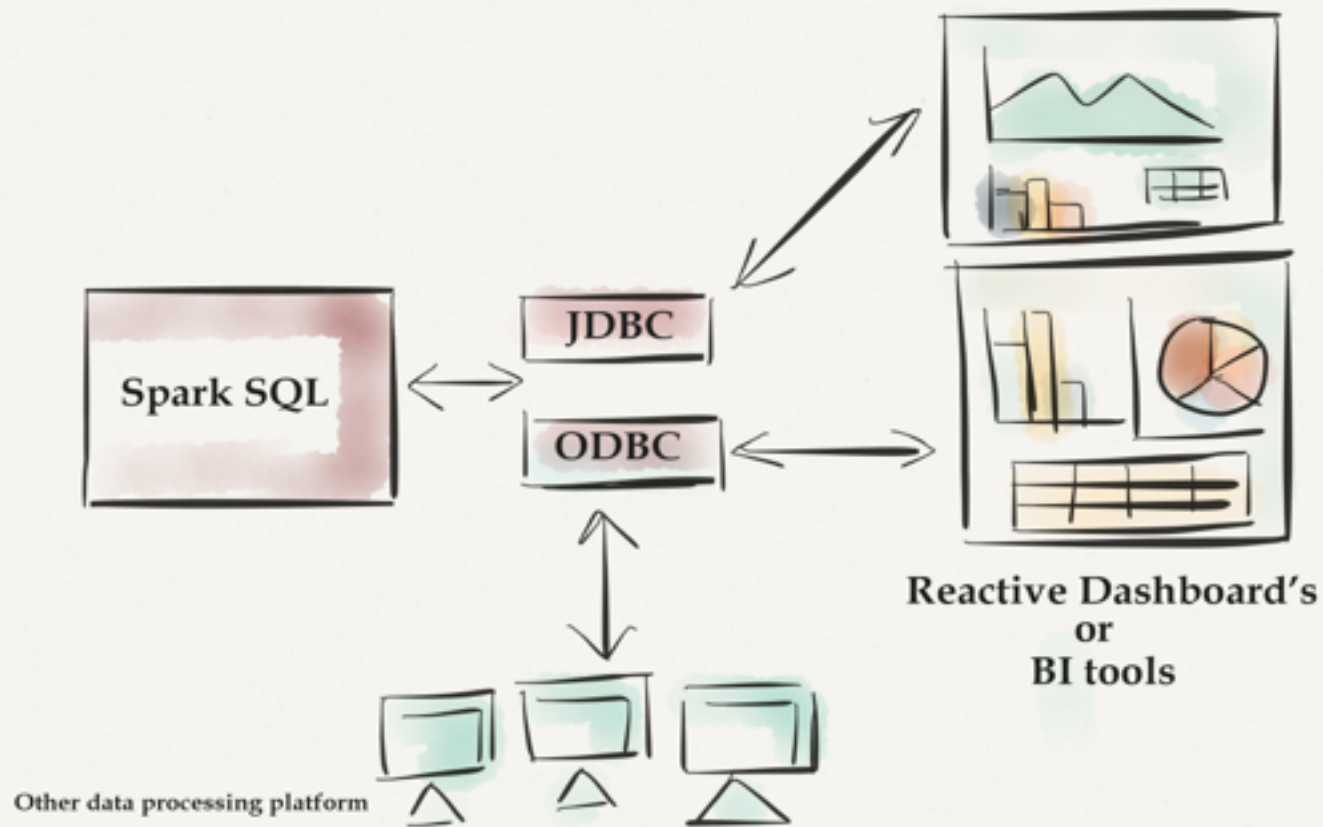
RDD Name	Storage Level	Cached Partitions	Fraction Cached	Size in Memory	Size in ExternalBlockStore	Size on Disk
MapPartitionsRDD	Memory Deserialized 1x Replicated	2	100%	3.3 KB	0.0 B	0.0 B

Spark SQL Introduction



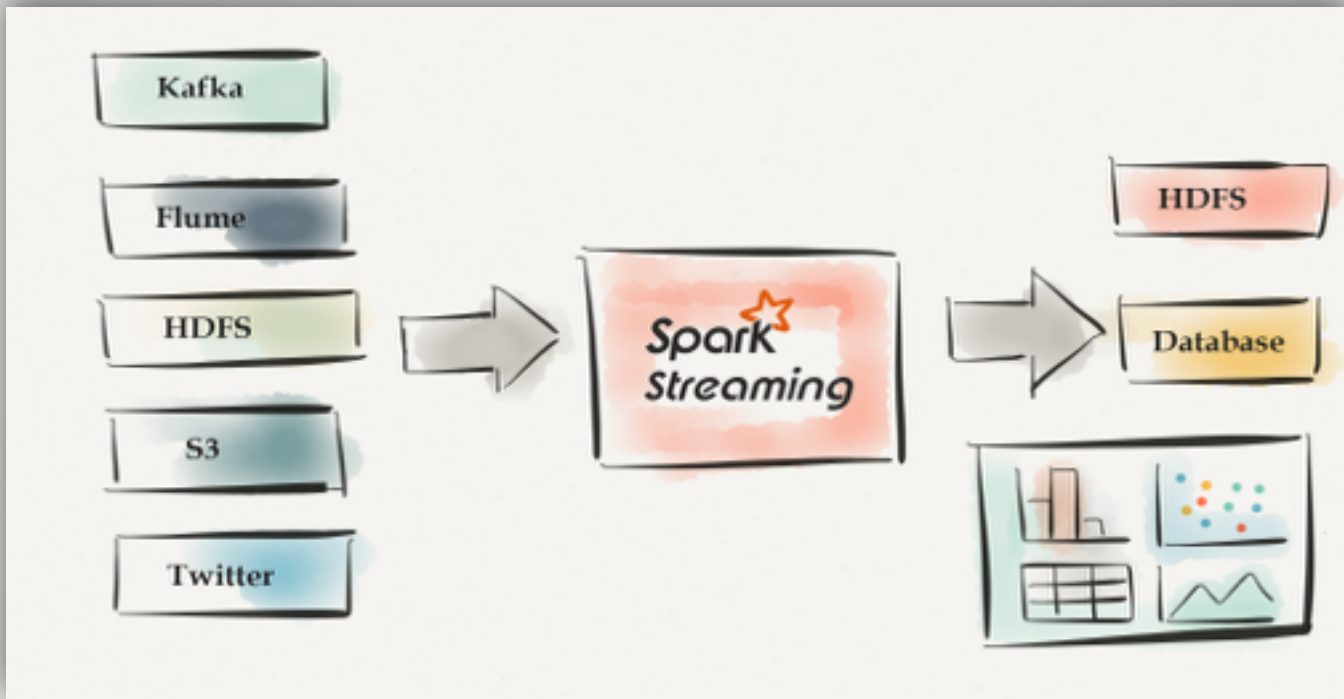
Spark SQL is Spark's module for working with structured data.


- Mix SQL queries with Spark programs
- Uniform Data Access, Connect to any data source
- DataFrames and SQL provide a common way to access a variety of data sources, including Hive, Avro, Parquet, ORC, JSON, and JDBC.
- Hive Compatibility Run unmodified Hive queries on existing data.
- Connect through JDBC or ODBC.



Spark Streaming Introduction

Spark Streaming is an extension of the core Spark API that enables **scalable, high-throughput, fault-tolerant stream processing** of live data streams.





\$git clone <https://github.com/rahulkumar-aws/WordCount.git>

\$ nc -lk 9999

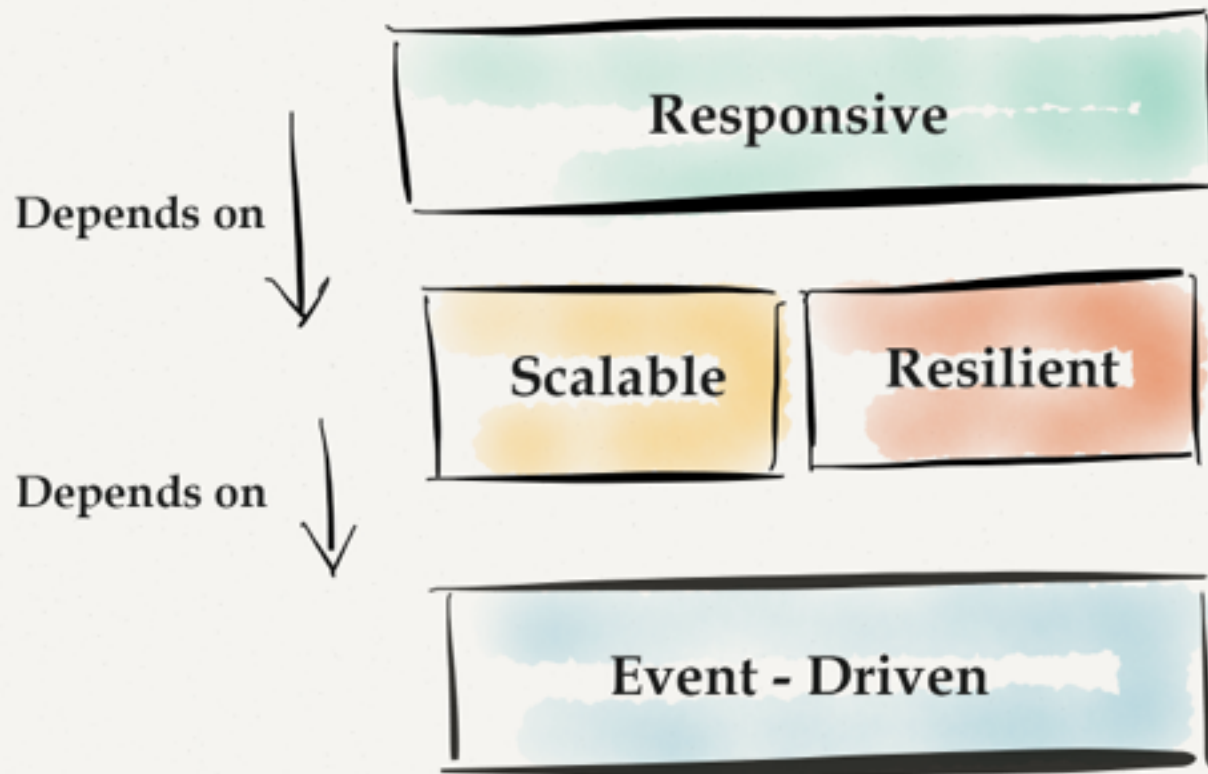
sbt "run-main StreamingWordCount"

Reactive Application



- Responsive
- Resilient
- Elastic
- Event Driven

<http://www.reactivemanifesto.org>



Typesafe Reactive Platform



Play Framework



The High Velocity Web Framework For Java and Scala

- RESTful by default
- JSON is a first class citizen
- Web sockets, Comet, EventSource
- Extensive NoSQL & Big Data Support

<https://www.playframework.com/download>

<https://downloads.typesafe.com/typesafe-activator/1.3.5/typesafe-activator-1.3.5-minimal.zip>

Akka is a toolkit and runtime for building highly concurrent, distributed, and resilient message-driven applications on the JVM.

- Simple Concurrency & Distribution
- Resilient by Design
- High Performance
- Elastic & Decentralized
- Extensible

Akka uses **Actor Model** that raise the abstraction level and provide a better platform to build **scalable**, **resilient** and **responsive applications**.

Demo

References

https://www.cs.berkeley.edu/~matei/papers/2012/nsdi_spark.pdf

<http://spark.apache.org/docs/latest/quick-start.html>

Learning Spark Lightning-Fast Big Data Analysis

By Holden Karau, Andy Konwinski, Patrick Wendell, Matei Zaharia

<https://www.playframework.com/documentation/2.4.x/Home>

<http://doc.akka.io/docs/akka/2.3.12/scala.html>

Thank You



Rahul Kumar rahul.k@sigmoid.com @rahul_kumar_aws